

兼具高效能與節能的 Edge/Device AI 晶片技術

ITRI deep Learning Accelerator design, system, tools, and applications

羅賢君 Shien-Chun Luo

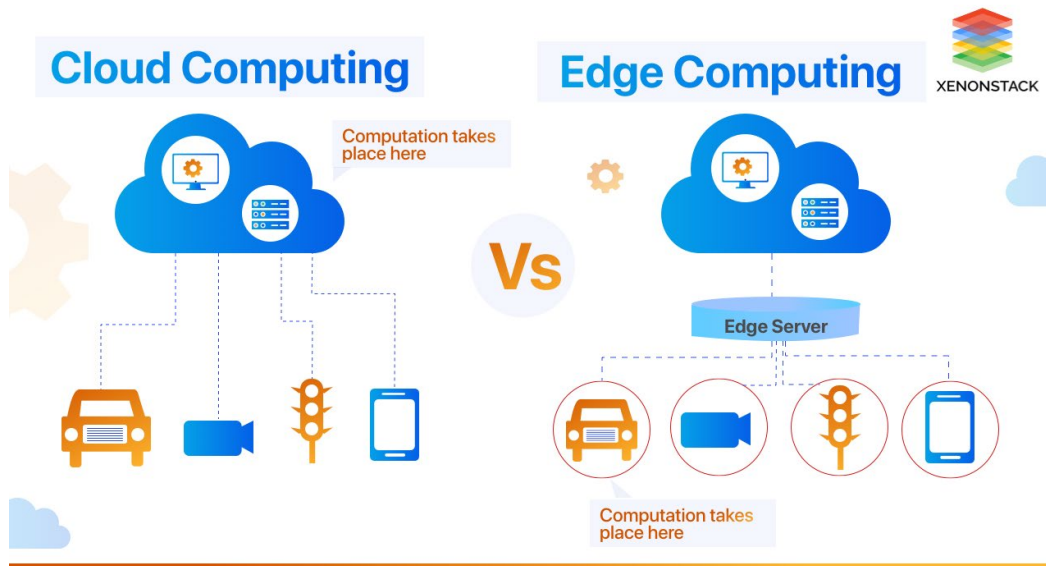
工業技術研究院 Industrial Technology Research Institute (ITRI)

電光系統研究所 Electronic and Optoelectronic System Research Lab (EOSL)

感知運算晶片系統組 Integrated Perception and Computation Systems (P)


What does AI Need ?

1. Massive **Data** acquire
2. Massive **Data** transfer
3. Massive **Data** compute



Why Need Edge AI ?

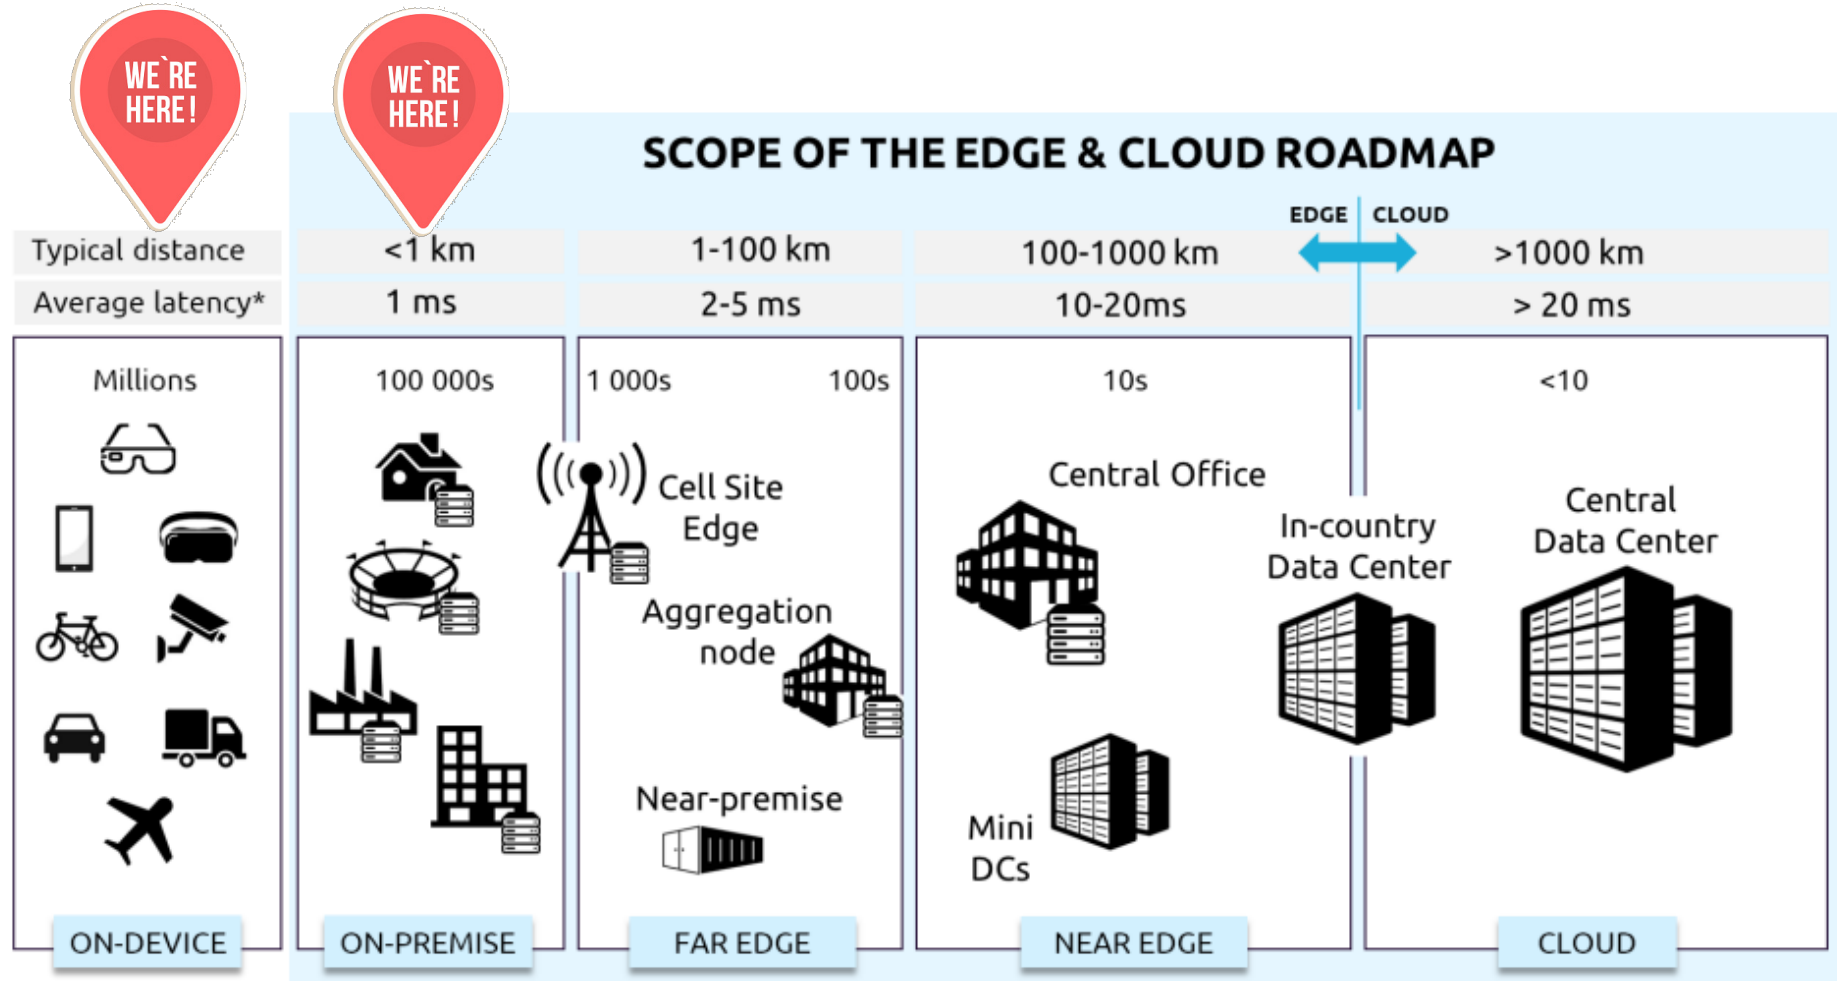
1. Latency
2. Privacy
3. Green



<https://www.workspot.com/blog/the-public-cloud-drive-to-lowest-latency-infrastructure/>

Where is Our (ITRI-AI-Chip) in the Edge?

1. Device 2. Edge Server



Source : European industrial technology roadmap for the next generation cloud-edge offering

Conventional and Deep-learning AI

Rule-based, conventional machine learning

Inspired by experience, observations, priority relations

Analytical solution, fast and accurate for simple questions

Too many rules and low accuracy for complex classification

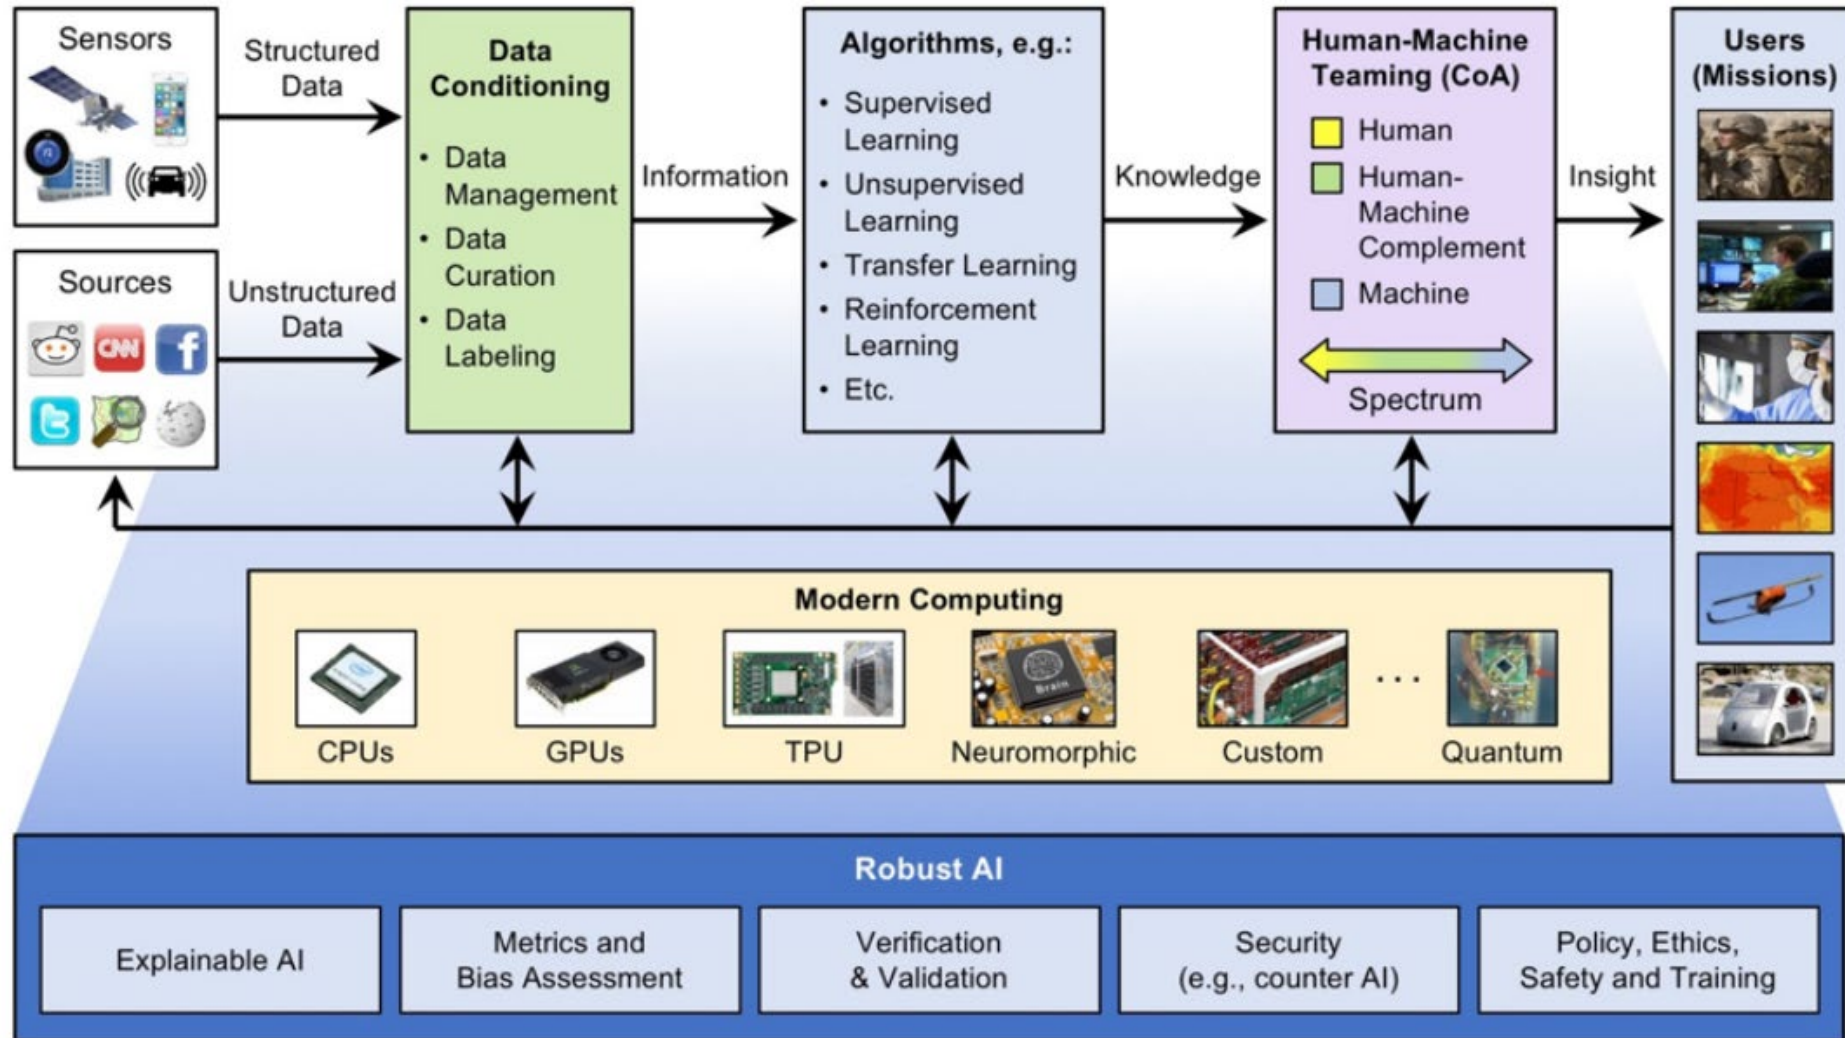
Deep-learning machine learning

Summarize and regress massive data, may be labelled consistently

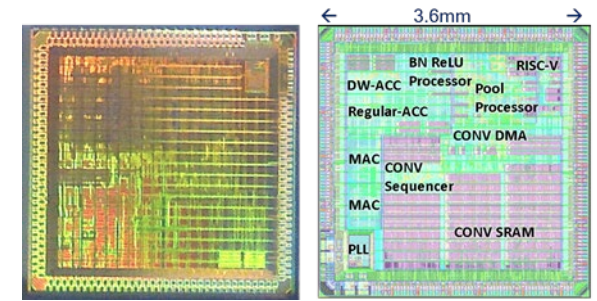
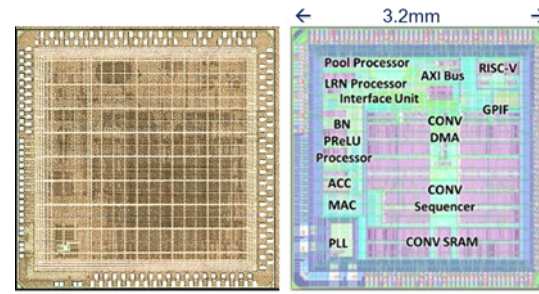
High-order, deep (> 3) relation among filters and inputs

Including supervised or non-supervised learning

AI, Machine Learning Architecture



AI Chip Products



AI accelerator IP or IC



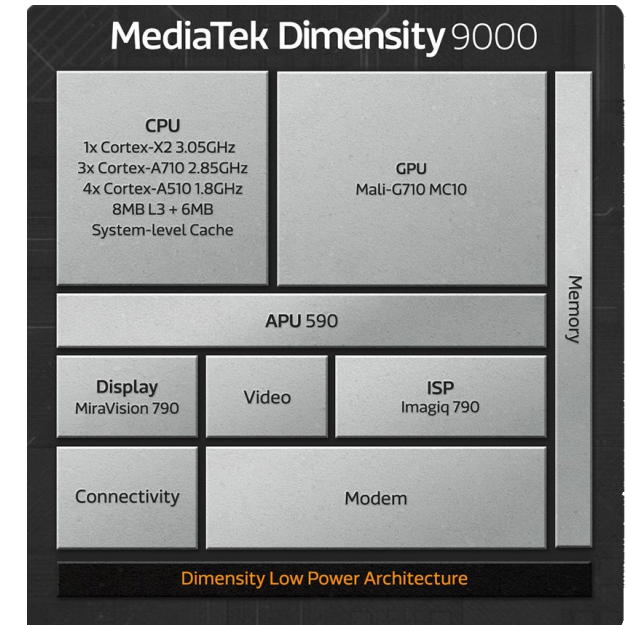
AI accelerator dongles or cards



AI embedded DEV boards

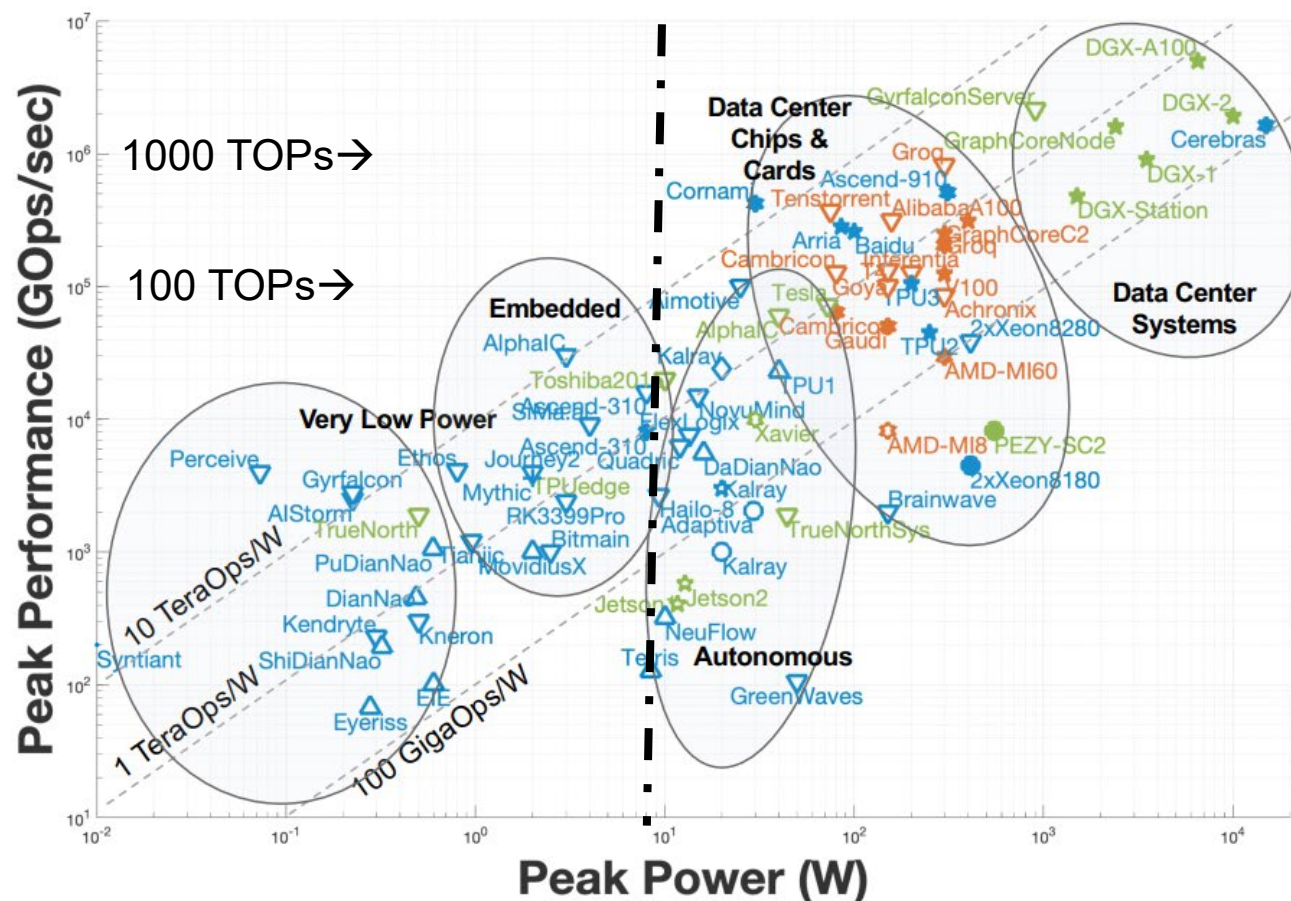


AI embedded ISP chip and camera module



AI embedded (APU, NPU...) SOC CPU AP IC

Performance Distribution of Machine Learning Hardware



Legend

Computation Precision

- + analog
- ◀ int1
- ▶ int2
- int4.8
- ▼ int8
- ◆ int8.32
- ▲ int16
- int12.16
- × int32
- ★ fp16
- ☆ fp16.32
- fp32
- * fp64

Form Factor

- Chip
- Card
- System

Computation Type

- Inference
- Training

Machine Learning HW

Performance Metrics

Performance index :

Tera-operators per second (TOPs)



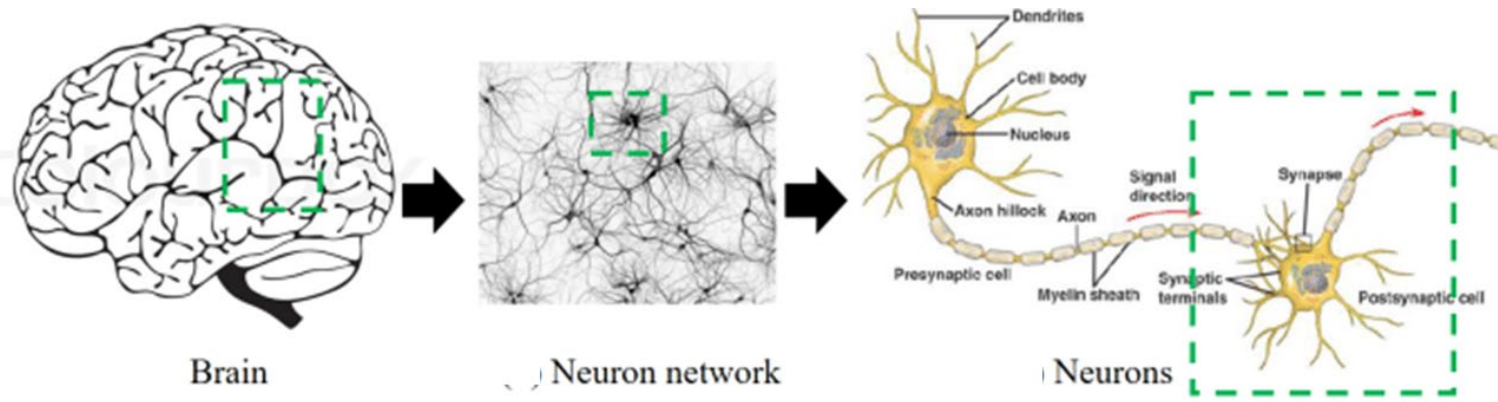
Energy efficiency index :

TOPs per Watt (TOPs/W)

AI Chip Glossary


- operators : add, multiply, shift, arithmetic functions
- precision : floating-point (FP16,32,64), integer (INT1,2,4,8,16)

From Neurons to Neural Networks




Digital

Convolutional Neural Network Solution



Digital

Dense Neural Network Solution




Analog

Mixed-mode

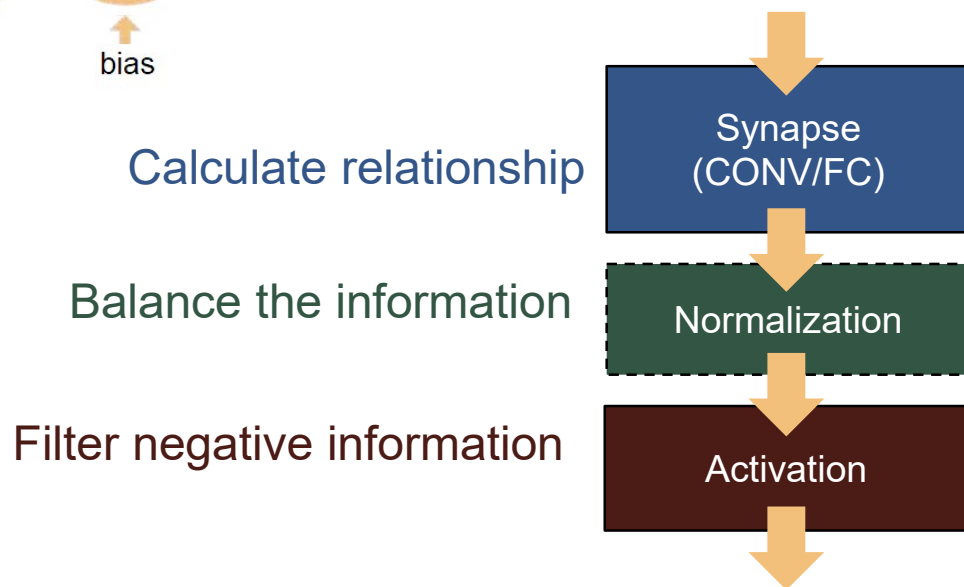
Neuromorphic, Processing in Memory

Neural Network Building Blocks and Graph

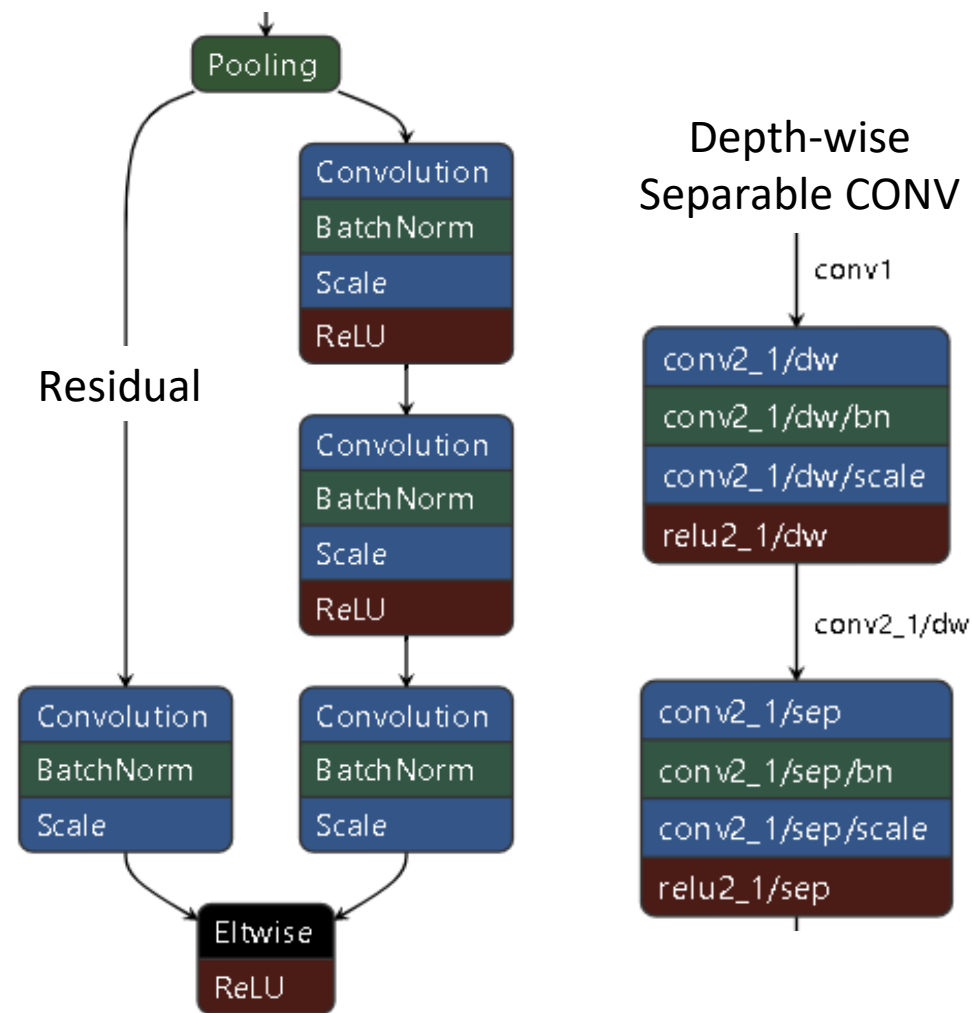
1. From Neurons to Functional Operators



2. From Operators to Functional Queues

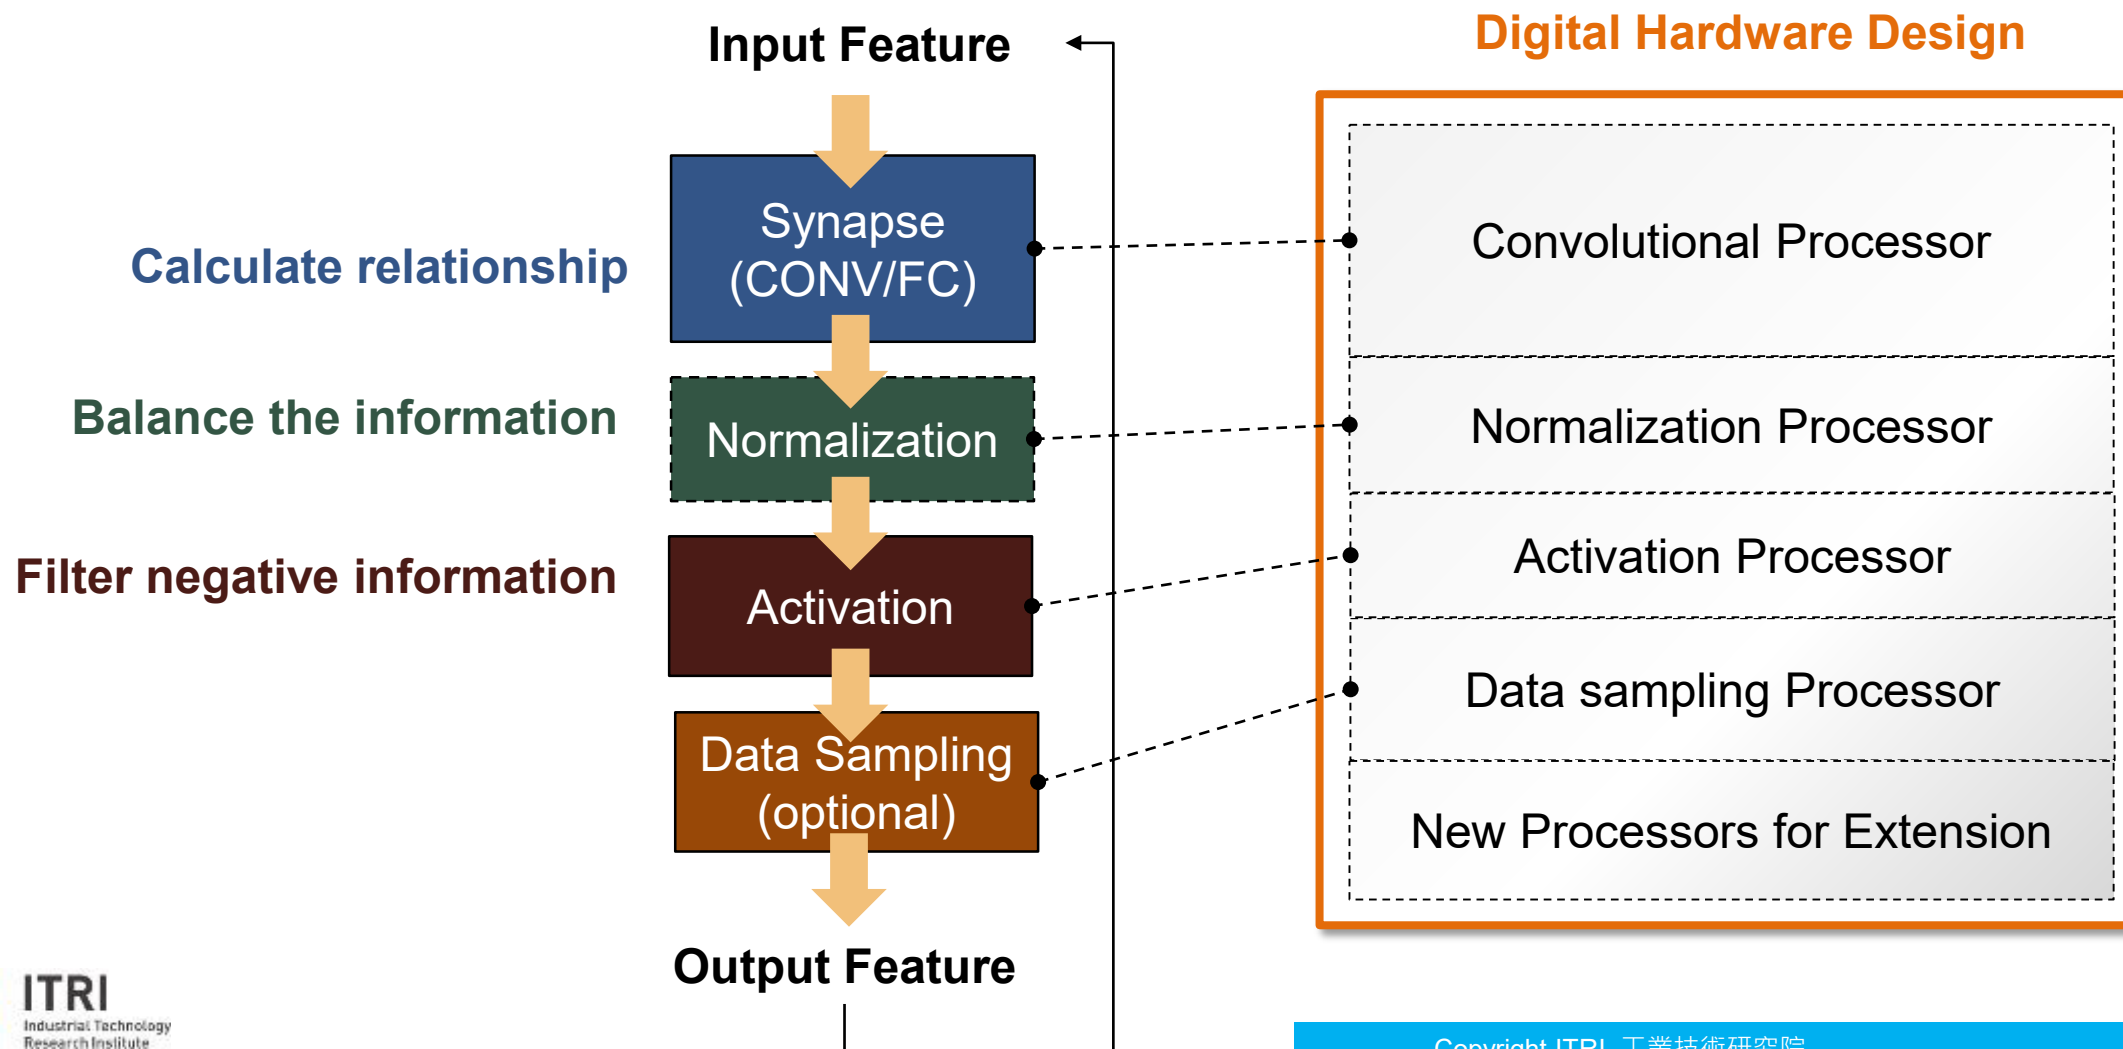


3. From Functional Queues to Neural Networks



This is How a DLA Architecture Formed

to construct a configurable HW queue that fits a generic neural operator queue



Let's See Some Backbone CNN Models

What are those layer ? What do you find ? How to imagine ?

| type | patch size/ stride | output size | depth | #1×1 | #3×3 reduce | #3×3 | #5×5 reduce | #5×5 | pool proj | params | ops |
|----------------|-----------------------|----------------|-------|------|----------------|------|----------------|------|--------------|--------|------|
| convolution | 7×7/2 | 112×112×64 | 1 | | | | | | | 2.7K | 34M |
| max pool | 3×3/2 | 56×56×64 | 0 | | | | | | | | |
| convolution | 3×3/1 | 56×56×192 | 2 | | 64 | 192 | | | | 112K | 360M |
| max pool | 3×3/2 | 28×28×192 | 0 | | | | | | | | |
| inception (3a) | | 28×28×256 | 2 | 64 | 96 | 128 | 16 | 32 | 32 | 159K | 128M |
| inception (3b) | | 28×28×480 | 2 | 128 | 128 | 192 | 32 | 96 | 64 | 380K | 304M |
| max pool | 3×3/2 | 14×14×480 | 0 | | | | | | | | |
| inception (4a) | | 14×14×512 | 2 | 192 | 96 | 208 | 16 | 48 | 64 | 364K | 73M |
| inception (4b) | | 14×14×512 | 2 | 160 | 112 | 224 | 24 | 64 | 64 | 437K | 88M |
| inception (4c) | | 14×14×512 | 2 | 128 | 128 | 256 | 24 | 64 | 64 | 463K | 100M |
| inception (4d) | | 14×14×528 | 2 | 112 | 144 | 288 | 32 | 64 | 64 | 580K | 119M |
| inception (4e) | | 14×14×832 | 2 | 256 | 160 | 320 | 32 | 128 | 128 | 840K | 170M |
| max pool | 3×3/2 | 7×7×832 | 0 | | | | | | | | |
| inception (5a) | | 7×7×832 | 2 | 256 | 160 | 320 | 32 | 128 | 128 | 1072K | 54M |
| inception (5b) | | 7×7×1024 | 2 | 384 | 192 | 384 | 48 | 128 | 128 | 1388K | 71M |
| avg pool | 7×7/1 | 1×1×1024 | 0 | | | | | | | | |
| dropout (40%) | | 1×1×1024 | 0 | | | | | | | | |
| linear | | 1×1×1000 | 1 | | | | | | | | |
| softmax | | 1×1×1000 | 0 | | | | | | | | |

Table 1: GoogLeNet incarnation o

↑ Inception

Resnet →

| layer name | output size | 18-layer | 34-layer | 50-layer | |
|------------|-------------|---|---|---|--|
| conv1 | 112×112 | 7×7, 64, stride 2 | | | |
| conv2_x | 56×56 | 3×3 max pool, stride 2 | | | |
| | | $\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$ | |
| conv3_x | 28×28 | $\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$ | $\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$ | |
| conv4_x | 14×14 | $\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$ | $\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$ |
| conv5_x | 7×7 | $\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$ | $\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ | $\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$ |
| | 1×1 | average pool, 1000-d fc, softmax | | | |
| FLOPs | | 1.8×10 ⁹ | 3.6×10 ⁹ | 3.8×10 ⁹ | 7.6×10 ⁹ |

| Table 1. MobileNet Body Architecture | | |
|--------------------------------------|---------------------|----------------|
| Type / Stride | Filter Shape | Input Size |
| Conv / s2 | 3 × 3 × 3 × 32 | 224 × 224 × 3 |
| Conv dw / s1 | 3 × 3 × 32 dw | 112 × 112 × 32 |
| Conv / s1 | 1 × 1 × 32 × 64 | 112 × 112 × 32 |
| Conv dw / s2 | 3 × 3 × 64 dw | 112 × 112 × 64 |
| Conv / s1 | 1 × 1 × 64 × 128 | 56 × 56 × 64 |
| Conv dw / s1 | 3 × 3 × 128 dw | 56 × 56 × 128 |
| Conv / s1 | 1 × 1 × 128 × 128 | 56 × 56 × 128 |
| Conv dw / s2 | 3 × 3 × 128 dw | 56 × 56 × 128 |
| Conv / s1 | 1 × 1 × 128 × 256 | 28 × 28 × 128 |
| Conv dw / s1 | 3 × 3 × 256 dw | 28 × 28 × 256 |
| Conv / s1 | 1 × 1 × 256 × 256 | 28 × 28 × 256 |
| Conv dw / s2 | 3 × 3 × 256 dw | 28 × 28 × 256 |
| Conv / s1 | 1 × 1 × 256 × 512 | 14 × 14 × 256 |
| 5× Conv dw / s1 | 3 × 3 × 512 dw | 14 × 14 × 512 |
| Conv / s1 | 1 × 1 × 512 × 512 | 14 × 14 × 512 |
| Conv dw / s2 | 3 × 3 × 512 dw | 14 × 14 × 512 |
| Conv / s1 | 1 × 1 × 512 × 1024 | 7 × 7 × 512 |
| Conv dw / s2 | 3 × 3 × 1024 dw | 7 × 7 × 1024 |
| Conv / s1 | 1 × 1 × 1024 × 1024 | 7 × 7 × 1024 |
| Avg Pool / s1 | Pool 7 × 7 | 7 × 7 × 1024 |
| FC / s1 | 1024 × 1000 | 1 × 1 × 1024 |
| Softmax / s1 | Classifier | 1 × 1 × 1000 |

↑ MobileNet

This is How Our DLA Architecture Construct

High Performance Techniques

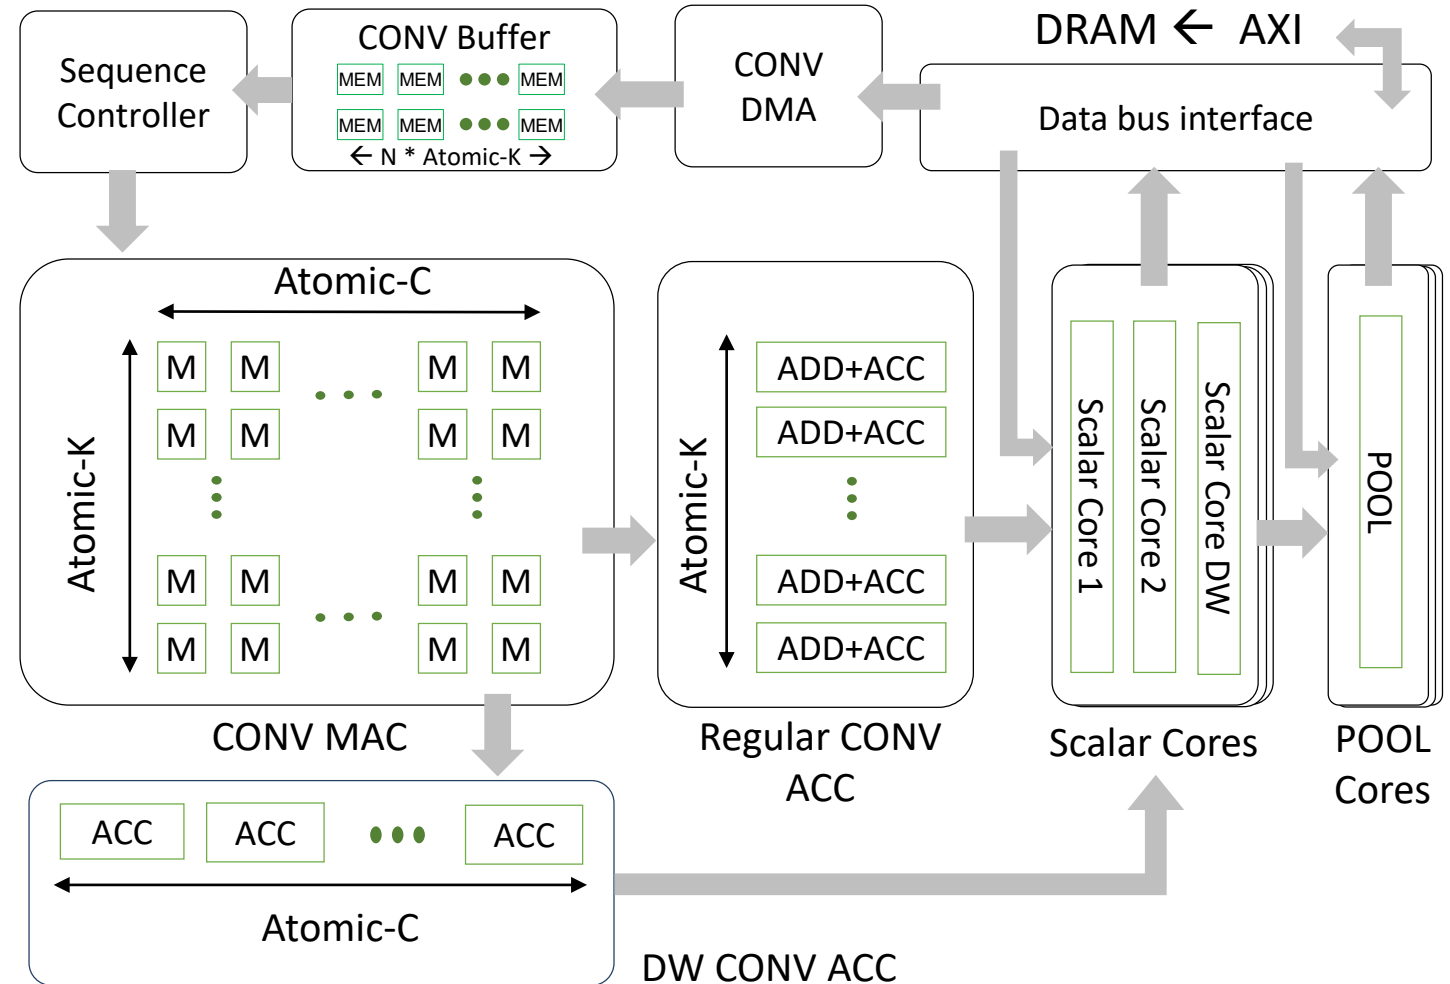
1. $2^C \times 2^K$ Multiplier array
2. Aligned CONV buffer array
3. Aligned accumulator array

Low-energy Techniques

1. Fused operators for saving data
2. Multiple scalars for OP swap
3. Identical IN & OUT data format

Ultra-low Power Techniques

1. Adaptive gating design
2. Data pipeline retiming








Optimize and Customize (1) : CNN Accelerator

Convolution contains many independent MACs and data overlap


1. Increase MAC PEs with high parallelism
2. Ensure the data supplement to those PEs
3. Improve energy efficiency, adaptive to the models

Concepts of step 1~3


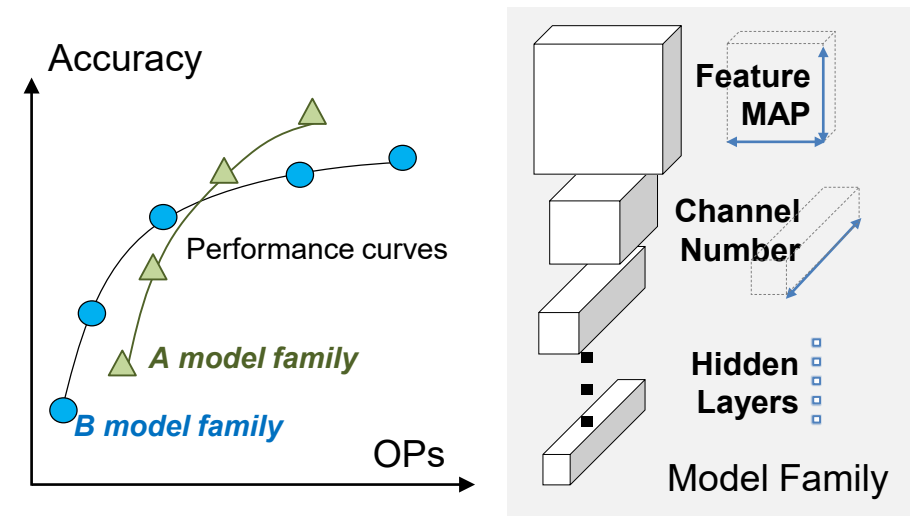


Optimize and Customize (2) : CNN Models


1. Select NN Model Backbone



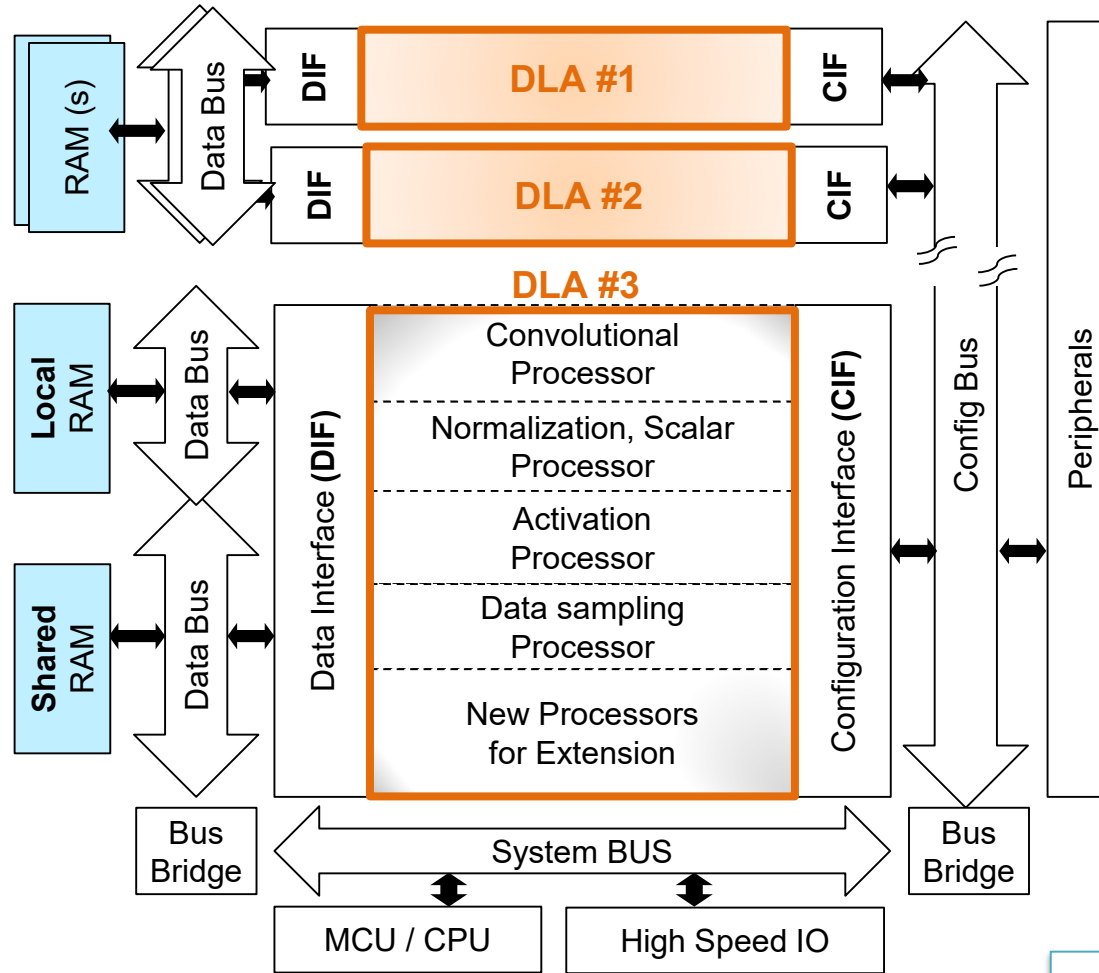
2. Search for the best configuration



3. Adapt to HW effective, HW specific features




ITRI Configurable and Scalable DLA Architecture

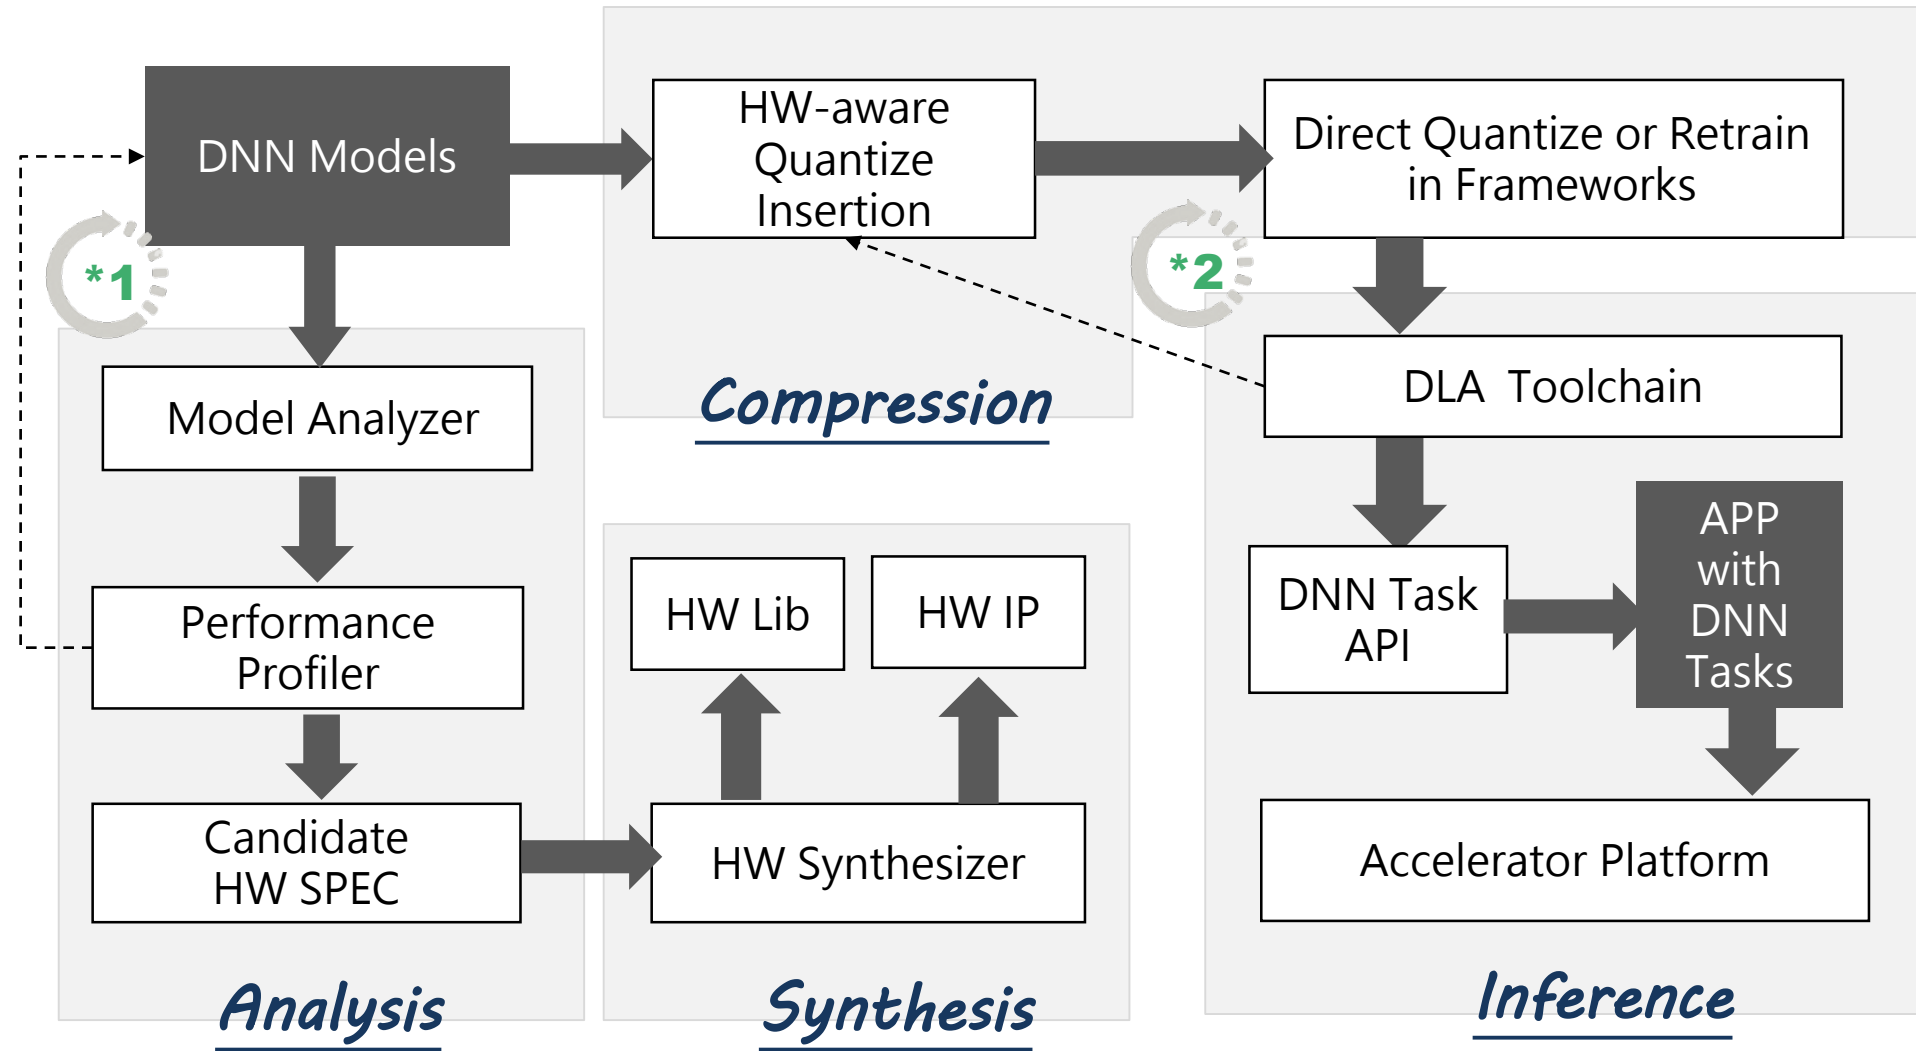


Our structure inherits NVDLA,
with various extensions

1. Configurable DLA number
2. Configurable MAC number
3. Configurable scalar cores for normalization, scale, quantize
4. Optional activation functions
5. Optional data sampling processors
6. Optional new processors
7. Custom CPU or MCU as the host



ITRI Service, Develop Flow of the DLA Platform



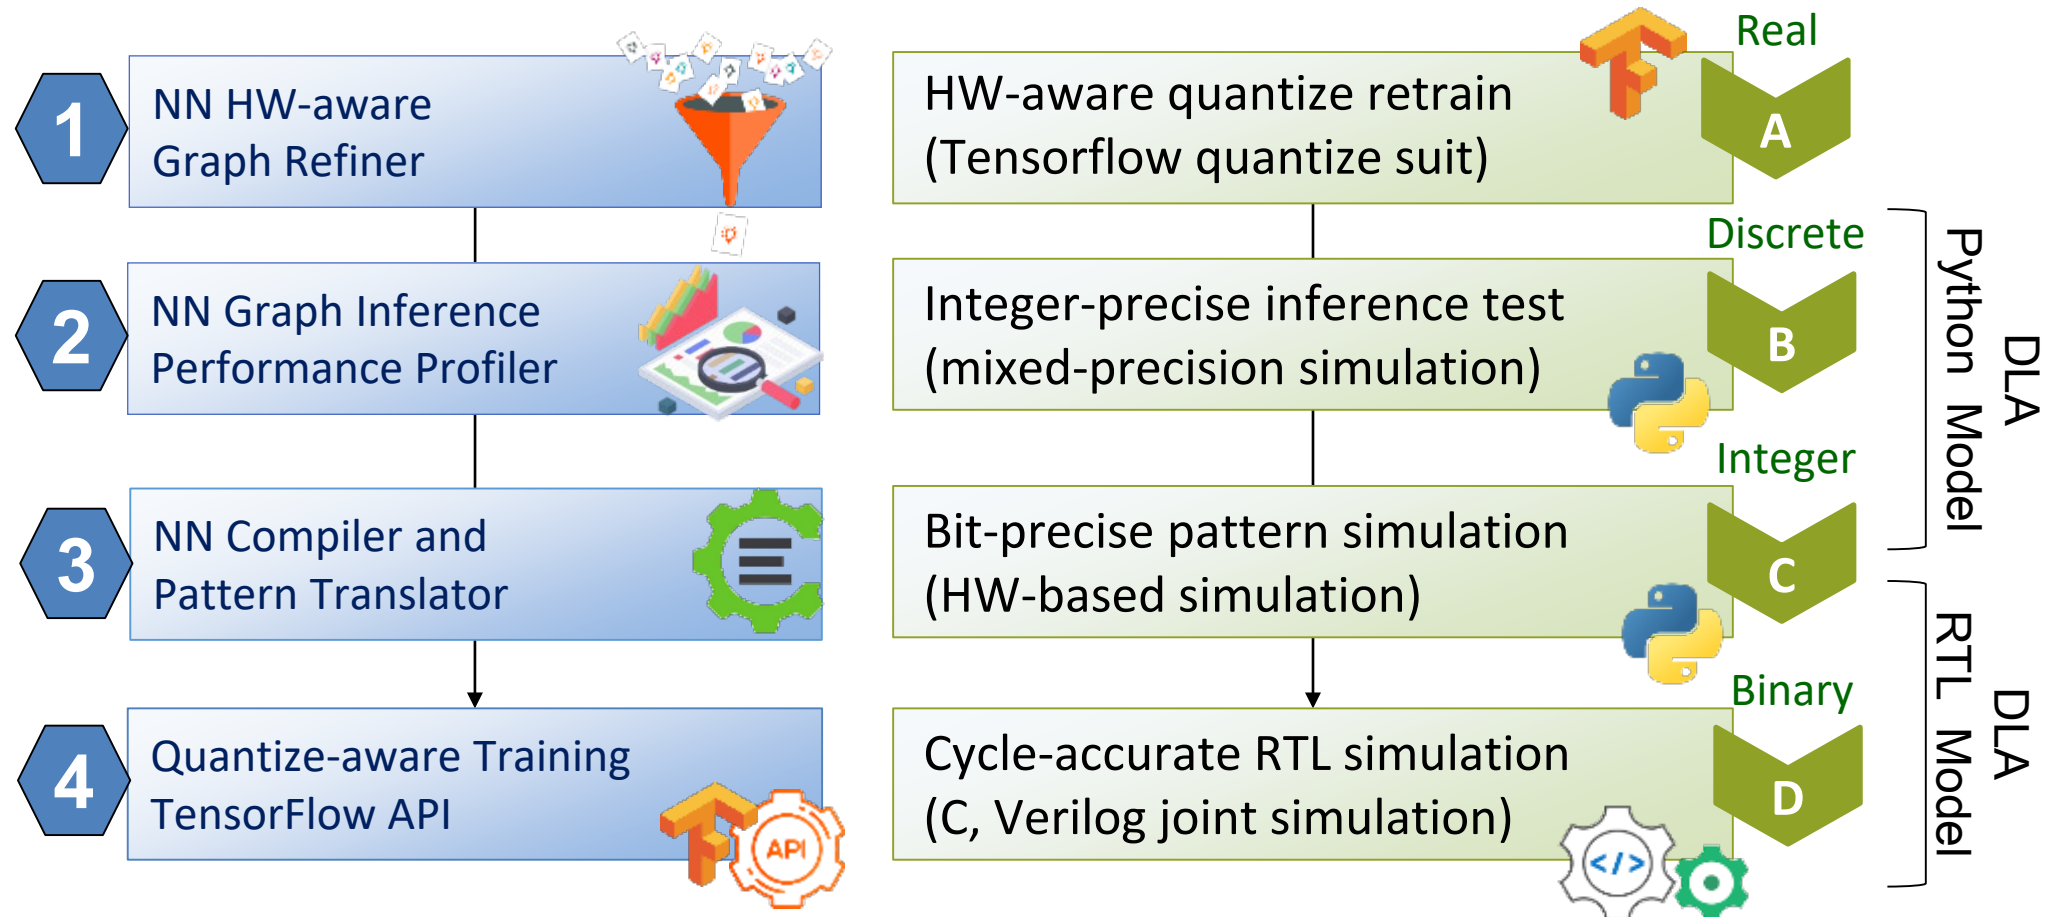
*1. Network architecture search

*2. HW-precise quantize modification

ITRI DLA Tool Chain and Verification Flow

from QAT, Profile, Bare-metal Compile, Simulate, to Deploy

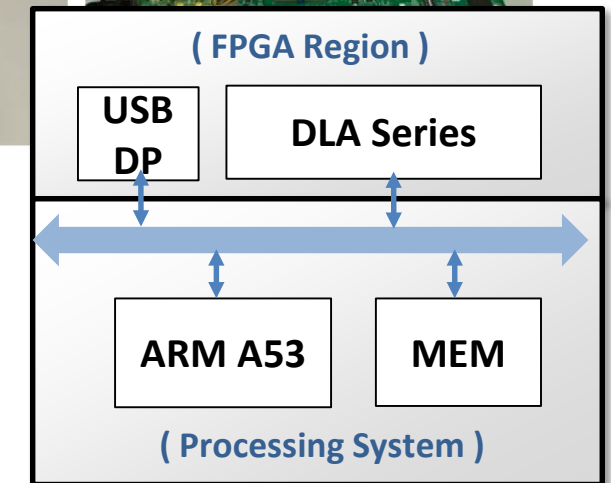


4 Main Tools + 4 Numerical Steps : linking from training to inference



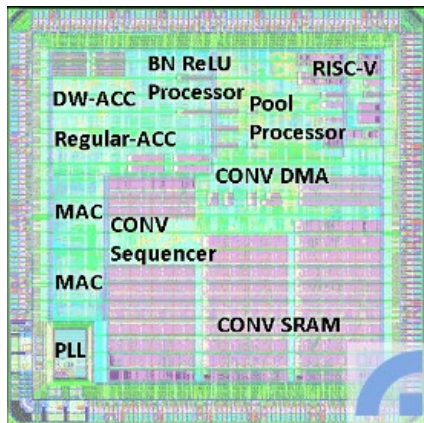

ITRI DLA IP and Standalone FPGA Prototype

| DLA Series | DLA256M | DLA2048M |
|------------------------------------|------------------|--------------------|
| MAC Number | 256 | 2048 |
| CONV SRAM | 128KB | 512KB |
| Norm. Engine | 1X | 8X |
| Pool Engine | 1X | 4X |
| AXI Data BW | 1X | 8X |
| FPGA Prototype | 43% Logic, 3%DSP | 78% Logic, 86% DSP |
| FPGA Frequency | 200 MHz | 125 MHz |
| Reference Model Speed (fps) | | |
| Tiny YOLO v1 | 14.4 | 28.7 |
| Tiny YOLO v2 | 7.5 | 28.7 |
| Tiny YOLO v3 | 12.3 | 32.3 |
| Full YOLO v3 | 1.07 | 4.5 |
| Full YOLO v4 | 1.04 | 4.4 |
| Resnet50 | 6.4 | 17.6 |
| Mobilenet | 33 | 110 |


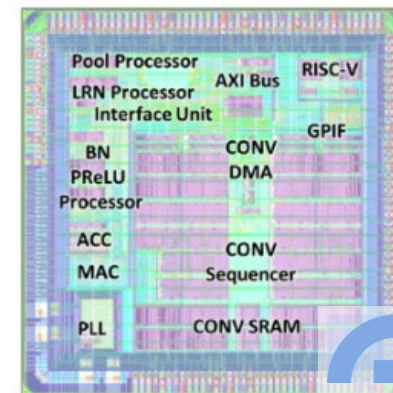

Standalone Xilinx ZCU102 FPGA Prototype



DLA + RISC-V as a USB Accelerator

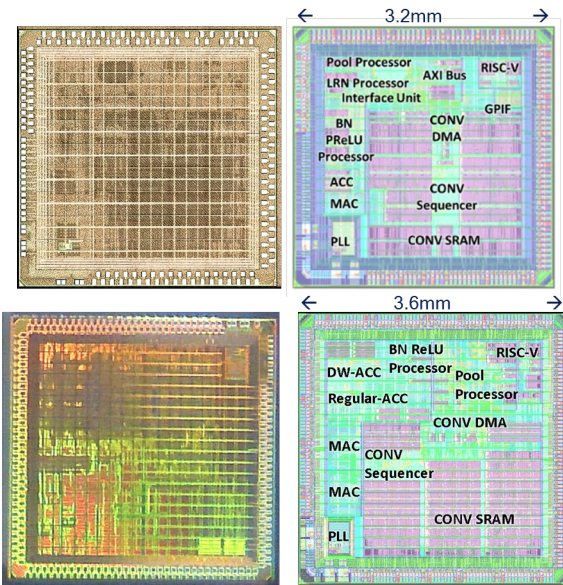


DLA 256M @65nm, 200 GOPs,
80 mW, 2.5 TOPs/W

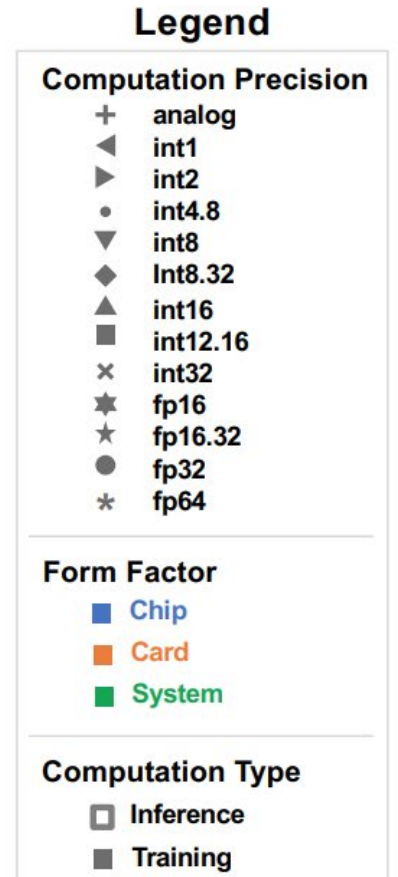
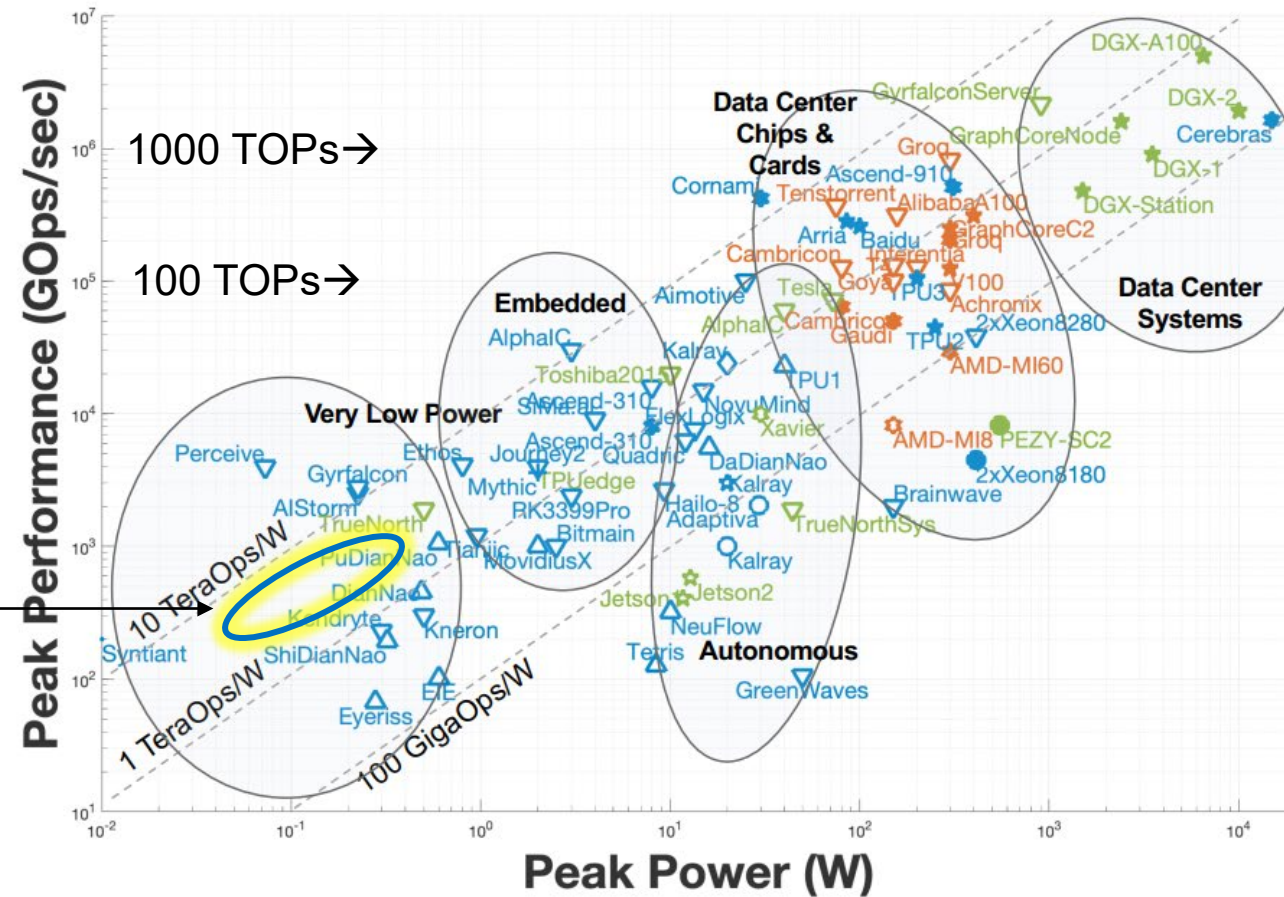


DLA 64M @65nm, 50 GOPs,
60 mW, 0.83 TOPs/W

ITRI DLA AI Chip Position on the Accelerator Map



256MAC @ 65nm,
200 GOPs/ 80mW,
2.5 TOPs/W



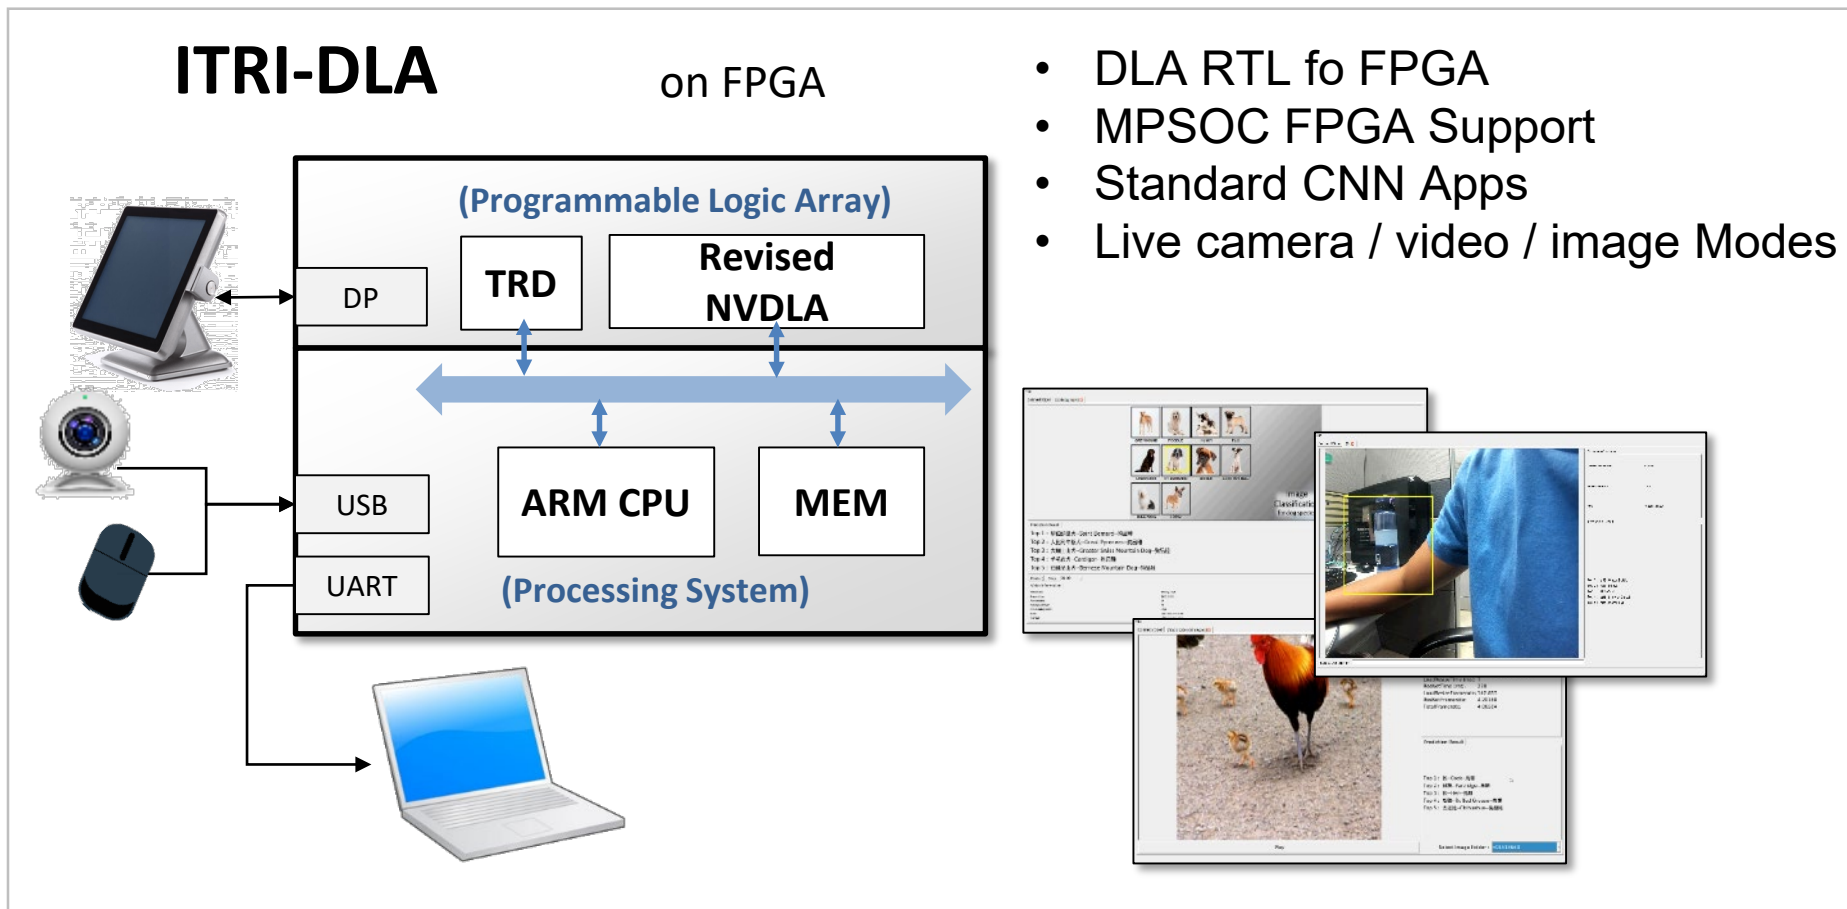
Summaries

1. Introduction of the need of edge and device AI.
2. Introduction of the ITRI DLA AI chip.
3. Like CPU/GPU/DSP, AI accelerator need its own SDK.
4. AI chip's SW works actually >> HW works.
5. AI chips can be embedded or plug-in dongles.
6. No heat sink nor fan makes AI edge and device fly.
7. In-situ AI decision saves massive raw data movement→ green.



ITRI-OpenDLA FPGA for Trial Evaluation

<https://github.com/SCLUO/ITRI-OpenDLA>



THANK YOU!

Questions and Comments?



ITRI-OpenDLA
<https://github.com/SCLUO/ITRI-OpenDLA>



DLA Perf Profiler
<https://github.com/SCLUO/Open-DLA-Performance-Profiler>